

## Contents

- [Contents](#)
- [Intel Pentium instruction-set specification](#)
  - [Opcodes](#)
  - [One-byte opcodes](#)
  - [Two-byte opcodes](#)
  - [Operands and effective addresses](#)
  - [Mod R/M opcodes](#)
  - [Operand-size and address-size prefixes](#)
  - [Floating-point opcodes](#)
  - [Arithmetic instructions](#)
  - [Other instructions \(in alphabetical order\)](#)
  - [Synthetic instructions](#)
  - [Assembly Specification for Gnu-Linux Assembler](#)
    - [Assembly names for opcodes](#)
    - [Assembly formats for operands](#)
    - [Assembly syntax for constructors](#)
  - [Miscellaneous](#)

## Intel Pentium instruction-set specification

This specification describes the Intel Pentium **[cite [intel:pentium](#)]**. At the instruction-set level, this specification could almost be used to describe the 486; the Pentium supports just a few instructions not found on the 486. This specification has *not* been used in an application, which means that it is probably full of bugs. **[We have plans for a test harness that should generate instructions at random, making sure that we get the same object code no matter whether we emit binary or assembly language, but as of July 1994 that harness is not in place.]**

Specifying the x86 is very different from specifying a RISC machine. In particular, it is difficult to know how to factor this architecture, or indeed whether to try to factor it at all. We've ended by trying to factor the opcodes into tables, but not to factor the instructions. Handling the x86 opcode tables is painful, but we prefer it to specifying the opcodes individually, because we believe the tables reduce the likelihood of error. Factoring the instructions turned out to be a hopeless exercise, with one exception: there is a family of 8 groups of arithmetic instructions that factor nicely.

The x86 is not quite sure whether it is an 8-bit, a 16-bit, or a 32-bit machine. We don't know the exact history, but we believe that the machine started life as an 8-bit machine, and that new opcodes were added when the architecture was extended to 16 bits. New opcodes were *not* added when it was extended again to 32 bits; instead, the presence or absence of a prefix is used to distinguish 16- from 32-bit instructions. Unfortunately, the encoding isn't completely specified at assembly time; the meaning of the prefix depends on the setting of a bit in the ``executable-segment descriptor." We have chosen to specify encodings in which instructions without prefixes operate on 32-bit quantities and the prefix selects 16-bit operation, but the encoding can be changed by changing just two lines in Section [\[->\]](#). An application that wanted to be able to switch between encodings dynamically would have to use the toolkit to generate two sets of encoding procedures, one each for the 16- and 32-bit defaults.

One thing we do is provide a way to generate just the 32-bit subset:

```
<pentium32.spec>= (U-> U->)
keep <32-bit constructors> <floating-point constructors>
```

```
<pentium-int.spec>=
discard <floating-point constructors>
```

---

---

```
height 1pt
[5.8em] --- 9One-Byte Opcode Map (page 0)
0 8 0|6ADD|PUSH!POP62
|Eb,Gb?Ev,Gv?Gb,Eb?Gv,Ev?AL,Ib?eAX,Iv|ES!ES |PUSH!POP62
|Eb,Gb?Ev,Gv?Gb,Eb?Gv,Ev?AL,Ib?eAX,Iv|SS!SS |SEG!DAA62
|Eb,Gb?Ev,Gv?Gb,Eb?Gv,Ev?AL,Ib?eAX,Iv| =ES! |SEG!AAA62
|Eb,Gb?Ev,Gv?Gb,Eb?Gv,Ev?AL,Ib?eAX,Iv| =SS! 4|8INC general register8
|eAX?eCX?eDX?eBX?eSP?eBP?eSI?eDI 8
|eAX?eCX?eDX?eBX?eSP?eBP?eSI?eDI 6|PUSH! POPA! BOUND!ARPL!
SEG!SEG!Operand!Address|PUSHAD!POPAD!Gv,Ma!Ew,Gw!=FS!=GS!Size! Size |JO?JNO?JB/JNAE/J?JNB/JAE/J?JZ?JNZ?JBE?JNBE
```

8|2Immediata Grp1!MOVb^\*!Grp1!2TEST!2XCHG224  
 |Eb,lb?Ev,lv!AL,immed!Ev,lb!Eb,Gb?Ev,Gv!Eb,Gb?Ev,Gv 9|NOP!7XCHG word to double-word register with eAX17  
 |!eCX?eDX?eBX?eSP?eBP?eSI?eDI A|4MOV|MOVSB!MOVsw!CMPSB!CMPsw44  
 |AL,Ob?eAX,Ov?Ob,AL?Ov,eAX|Xb,Yb!Xv,Yv!Xb,Yb!Xv,Yv B|8MOV immediate byte into register8  
 |AL?CL?DL?BL?AH?CH?DH?BH C|2Shift Grp2a!2RET near|LES!LDS!2MOV422  
 |Eb,lb?Ev,lb!lw? |Gv,Mp?Gv,Mp!Eb,lb?Ev,lv |AAM!AAD!\*!XLAT44  
 |Eb,1?Ev,1?Eb,CL?Ev,CL|!!! E|LOOPNE!LOOPE!LOOP!JCXZ/JEC!2IN!2OUT44  
 |jb!jb!jb!jb!AL,lb?AX,lb!lb,AL?lb,eAX 62  
 |!!!REPE!!!Eb?Ev **Pentium opcodes (page 0)[\*]**

---

height 1pt  
 [5.8em] --- **9One-Byte Opcode Map (page 1)**  
 816 0|6OR|PUSH!2-byte62  
 |Eb,Gb?Ev,Gv?Gb,Eb?Gv,Ev?AL,lb?eAX,lv|CS!escape |PUSH!POP62  
 |Eb,Gb?Ev,Gv?Gb,Eb?Gv,Ev?AL,lb?eAX,lv|DS!DS |SEG!DAS62  
 |Eb,Gb?Ev,Gv?Gb,Eb?Gv,Ev?AL,lb?eAX,lv| =CS! |SEG!AAS62  
 |Eb,Gb?Ev,Gv?Gb,Eb?Gv,Ev?AL,lb?eAX,lv| =DS! 4|8DEC general register8  
 |eAX?eCX?eDX?eBX?eSP?eBP?eSI?eDI 8  
 |eAX?eCX?eDX?eBX?eSP?eBP?eSI?eDI 6|PUSH!MUL! PUSH!MUL! INSB! INSW/D!OUTSB!OUTSW/D|lv! Gv,Ev,lv!lb!  
 Gv,Ev,lb!Yb,DX!Yv,DX! DX,Xb!DX,Xv |JS?JNS?JP?JNP?JL?JNL?JLE?JNLE 8|4MOV|MOV!LEA!MOV!POP44  
 |Eb,Gb?Ev,Gv?Gb,Eb?Gv,Ev|Ew,Sw!Gv,M!Sw,Ew!Ev 9|CBW!CWD/CDQ!CALL!WAIT|PUSHF!POPF!SAHF!LAHF|!aP!|Fv!Fv!!  
 !STOSB!STOSW/D|LODSB!LODSW/D!SCASB!SCASW/D26  
 |AL,lb?eAX,lv!Yb,AL!Yv,eAX|AL,Xb!eAX,Xv!AL,Yb!eAX,Yv B|8MOV immediate word or double into word or double  
 register|eAX!eCX!eDX!eBX!eSP!eBP!eSI!eDI C|ENTER!LEAVE!RET far!RET far!INT!INT!INTO!IRET|lw,lb!!lw!!3!lb!! D|8ESC  
 (Escape to coprocessor instruction set)|8 E|CALL!3|MP!2IN!2OUT17  
 |jv!jv?Ap?jb!AL,DX?eAX,DX!DX,AL?DX,eAX **Pentium opcodes (page 1)[\*]**

Intel uses some naming conventions to try to tame the confusion surrounding operands. ``Operand specifiers" describe the locations and sizes of operands. The specifiers are (mostly) composed from the following pieces:

Operand	Width
E effective address (memory or register)	b 8-bit bytes w 16-bit words
G general-purpose register	d 32-bit doublewords
I immediate	v variable (w or d)

For example, the specifier ``Eb,Gb" describes an 8-bit memory-to-register instruction. ``Ev,Gv" describes a similar instruction that operates on 16 or 32 bits, depending on the presence or absence of a prefix. Most of the x86 instructions are overloaded, but the opcode tables often use operand specifiers as suffixes to distinguish them. In some cases, where the machine specification doesn't give distinguishing suffixes, we have invented some.

Many of the instructions support all three sizes, and we specify them in two variants: a b variant with no prefix, and v variants with and without prefixes. When the two v variants differ only in the presence or absence of a prefix, we can specify them simultaneously using the ov pattern, which we define in Section [\[->\]](#) to mean ``optional prefix." Sometimes, however, we have to specify all three variants explicitly, as when there is an immediate operand---in that case, we have to give three different output patterns because the token holding the immediate operand may be 8, 16, or 32 bits wide.

Here is the overall structure of the Intel specification:

[<pentium-core.spec>=](#)  
[<field specs>](#)  
[patterns <patterns for integer opcodes>](#)  
[<pattern specs for other patterns>](#)  
[<prefix assignments>](#)  
[<placeholders>](#)  
 relocatable reloc  
[<constructors for displacements>](#)  
[<constructors for effective addresses>](#)  
[<arithmetic constructors>](#)

constructors  
[<alphabetical constructors>](#)

## Opcodes

For other machines, we were able to specify entire opcode tables in single declarations. The Intel tables are less tractable, because they don't just contain opcodes; they contain a mix of opcodes and suffixes. We've broken most of the tables into pieces, as shown below. There's an argument for abandoning opcode tables entirely, using only constructors to describe the encodings, but we've decided to keep opcodes. This style of specification lets us use a little factoring, and it gives us a little protection against errors in the distributed opcodes, like the groupx opcodes.

### One-byte opcodes

We want to take advantage of factoring when possible, but it works well only for the ``arithmetic group," shown in the upper left corners of Figures [\[<-\]](#) and [\[<-\]](#), which represent the ``one-byte opcode map" from pages A-5 and A-6 of the Intel manual [[cite intel:pentium](#)]. This corner of the opcode table clearly represents an outer product of 8 opcodes with 6 suffixes, and we treat it as such.

Most of the rest of the opcode table we treat in purely geometric fashion, using row, column, and page numbers as Cartesian coordinates to determine opcodes.

[<field specs>=](#) ([<-U U->](#)) [[D->](#)]  
 fields of opcodet (8) row 4:7 col 0:2 page 3:3  
[<more fields of opcode>](#)

[<placeholders>=](#) ([<-U U->](#)) [[D->](#)]  
 placeholder for opcodet is HLT

Because the map is so chaotic, we break it into rows, for the most part specifying one or two rows at a time, but sometimes breaking rows into pieces. The first rows are among the most interesting; rows 0--3 contain the outer product of arithmetic operators with operand specifiers:

[<patterns for integer opcodes>=](#) ([<-U U->](#)) [[D->](#)]  
 arith is any of [ ADD OR  
                   ADC SBB  
                   AND SUB  
                   XOR CMP ], which is row = {0 to 3} & page = [0 1]  
 [ Eb.Gb Ev.Gv Gb.Eb Gv.Ev AL.Ib eAX.Iv ] is col = {0 to 5}

The other columns in rows 0--3 follow a less discernible pattern.

[<patterns for integer opcodes>+=](#) ([<-U U->](#)) [[<-D->](#)]  
 [ PUSH.ES POP.ES    PUSH.CS esc2  
   PUSH.SS POP.SS    PUSH.DS POP.DS  
   SEG.ES DAA        SEG.CS DAS  
   SEG.SS AAA        SEG.DS AAS ] is row = {0 to 3} & page = [0 1] & col = [6 7]

Rows 4 and 5 are the general-register opcodes, formed by an outer product of operation and register specifier. **[\*]** Although the register specifier is actually part of the opcode, we treat it as an operand below. We create the field r32 as an alias for col, so we can use special register names for the values.

[<patterns for integer opcodes>+=](#) ([<-U U->](#)) [[<-D->](#)]  
 regops is any of [ INC DEC  
                   PUSH POP ], which is row = [4 5] & page = [0 1]

[<more fields of opcode>=](#) ([<-U](#)) [[D->](#)]  
 r32 0:2

[<field specs>+=](#) ([<-U U->](#)) [[<-D->](#)]  
 fieldinfo r32 is [names [ eAX eCX eDX eBX eSP eBP eSI eDI ]]

[<more fields of opcode>+=](#) ([<-U](#)) [[<-D->](#)]  
 sr16 0:2

[<field specs>+=](#) ([<-U U->](#)) [[<-D->](#)]  
 fieldinfo sr16 is [sparse [ cs=1, ss=2, ds=3, es=4, fs=5, gs=6 ] ]

<more fields of opcode>+ = (<-U>) [<-D->]

r16 0:2

<field specs>+ = (<-U U->) [<-D->]

fieldinfo r16 is [names [ AX CX DX BX SP BP SI DI ]]

Row 6:

<patterns for integer opcodes>+ = (<-U U->) [<-D->]

[ PUSHB POPA BOUND ARPL SEG.FS SEF.GS OpPrefix AddrPrefix  
 PUSH.Iv IMUL.Iv PUSH.Ib IMUL.Ib INSB INSv OUTSB OUTSv  
 ] is page = [0 1] & row = 6 & col = {0 to 7}

Row 7 is factored so the jump codes can be re-used in the two-byte opcode map. Again, the row contains an outer product, but this time the columns determine jump conditions, not operand specifiers.

<patterns for integer opcodes>+ = (<-U U->) [<-D->]

Jb is row = 7

cond is any of [ .0 .NO .B .NB .Z .NZ .BE .NBE .S .NS .P .NP .L .NL .LE .NLE ],  
 which is page = [0 1] & col = {0 to 7}

Row 8 is a bit hard to follow because it contains a seemingly random mix of opcodes and operand specifiers. On page 0, we've left the "immediate Group1" implicit, giving only the operand specifiers. On page 1, MOV shares the operand specifiers we gave with the arithmetic instructions in rows 0--3.

<patterns for integer opcodes>+ = (<-U U->) [<-D->]

[ Eb.Ib Ev.Iv MOV.B Eb.Ib TEST.Eb.Gb TEST.Ev.Gv XCHG.Eb.Gb XCHG.Ev.Gv ] is  
 row = 8 & page = 0 & col = {0 to 7}

MOV is row = 8 & page = 1

[ MOV.Ew.Sw LEA MOV.Sw.Ew POP.Ev ] is row = 8 & page = 1 & col = {4 to 7}

On page 0, row 9 is XCHG (or NOP). Again, the register operand is actually part of the opcode. Page 1 has several opcodes.

<patterns for integer opcodes>+ = (<-U U->) [<-D->]

XCHG is row = 9 & page = 0

NOP is XCHG & col = 0

[ CBW CWDQ CALL.aP WAIT PUSHF POPF SAHF LAHF ] is row = 9 & page = 1 & col = {0 to 7}

Although there is an outer product or two lurking in row 10 (A), we don't try to specify it, because the operand specifiers used there aren't widely useful.

<patterns for integer opcodes>+ = (<-U U->) [<-D->]

[ MOV.AL.0b MOV.eAX.0v MOV.0b.AL MOV.0v.eAX MOVSB MOVSV CMPSB CMPSv  
 TEST.AL.Ib TEST.eAX.Iv STOSB STOSv LODSB LODSV SCASB SCASv  
 ] is row = 10 & page = [0 1] & col = {0 to 7}

Row 11 (B) is another row in which the register operand is implicit in the opcode, but we need to define a new field to denote 8-bit registers.

<patterns for integer opcodes>+ = (<-U U->) [<-D->]

MOVib is row = 11 & page = 0

MOViv is row = 11 & page = 1

<more fields of opcode>+ = (<-U>) [<-D>]

r8 0:2

<field specs>+ = (<-U U->) [<-D->]

fieldinfo r8 is [names [ AL CL DL BL AH CH DH BH ]]

Rows 12 and 13 contain the bit operators; the others are easy.

<patterns for integer opcodes>+ = (<-U U->) [<-D->]

[ B.Eb.Ib B.Ev.Ib RET.Iw RET LES LDS MOV.Eb.Ib MOV.Ev.Iv  
 B.Eb.1 B.Ev.1 B.Eb.CL B.Ev.CL AAM AAD \_ XLAT  
 ] is row = [12 13] & page = 0 & col = {0 to 7}

[ ENTER LEAVE RET.far.Iw RET.far INT3 INT.Ib INTO IRET ]

is row = 12 & page = 1 & col = {0 to 7}

ESC is row = 13 & page = 1

Neither of the last two rows can use factoring.

```
<patterns for integer opcodes>+= (<-U U->) [<-D->]
[ LOOPNE LOOPE LOOP JCXZ IN.AL.Ib IN.eAX.Ib OUT.Ib.AL OUT.Ib.eAX
  LOCK _ REPNE REP HLT CMC grp3.Eb grp3.Ev

  CALL.Jv JMP.Jv JMP.Ap JMP.Jb IN.AL.DX IN.eAX.DX OUT.DX.AL OUT.DX.eAX
  CLC STC CLI STI CLD STD grp4 grp5
] is page = [0 1] & row = [14 15] & col = {0 to 7}
```

## Two-byte opcodes

The two-byte tables are fairly sparse, so we haven't bothered to reproduce them in a table for this report. We describe them one page at a time beginning with page 0. The first token of each two-token pattern contains the esc2 opcode.

The first two rows are:

```
<patterns for integer opcodes>+= (<-U U->) [<-D->]
[ grp6 grp7 LAR LSL
  MOV.Eb.Gb MOV.Gv.Ev MOV.Gb.Eb MOV.Ev.Gv ]
is esc2; page = 0 & row = [0 1] & col = {0 to 3}
CLTS is esc2; page = 0 & row = 0 & col = 6
```

Row 3 is a block of MOV instructions.

```
<patterns for integer opcodes>+= (<-U U->) [<-D->]
[ MOV.Rd.Cd MOV.Rd.Dd MOV.Cd.Rd MOV.Dd.Rd MOV.Rd.Td _ MOV.Td.Rd ]
is esc2; page = 0 & row = 3 & col = {0 to 6}
```

Row 4:

```
<patterns for integer opcodes>+= (<-U U->) [<-D->]
[ WRMSR RDTSC RDMSR ] is esc2; page = 0 & row = 4 & col = {0 to 2}
```

The jumps in row 8 and sets in row 9 span two pages. They are outer products of opcodes and the conditions defined in row 7 of the one-byte opcode map.

```
<patterns for integer opcodes>+= (<-U U->) [<-D->]
Jv is esc2; row = 8
SETb is esc2; row = 9
```

Rows 10 and 11 are more madness:

```
<patterns for integer opcodes>+= (<-U U->) [<-D->]
[ PUSH.FS POP.FS CPUID BT SHLD.Ib SHLD.CL
  CMPXCHG.Eb.Gb CMPXCHG.Ev.Gv LSS BTR LFS LGS MOVZX.Gv.Eb MOVZX.Gv.Ew ]
is esc2; page = 0 & row = [10 11] & col = {0 to 7}
```

Row 12 is the only remaining non-empty row on page 0.

```
<patterns for integer opcodes>+= (<-U U->) [<-D->]
[ XADD.Eb.Gb XADD.Ev.Gv grp9 ] is esc2; page = 0 & row = 12 & col = [0 1 7]
```

There are even fewer opcodes on page 1.

```
<patterns for integer opcodes>+= (<-U U->) [<-D->]
[ INVD WBINVD ] is esc2; row = 0 & page = 1 & col = [0 1]
```

Rows 1--7 are empty, and 8 and 9 were covered on the previous page. Row 10 is:

```
<patterns for integer opcodes>+= (<-U U->) [<-D->]
[ PUSH.GS POP.GS RSM BTS SHRD.Ib SHRD.CL _ IMUL.Gv.Ev ]
is esc2; row = 10 & page = 1 & col = {0 to 7}
```

Row 11:

```
<patterns for integer opcodes>+= (<-U U->) [<-D->]
[ grp8 BTC BSR MOVXSX.Gv.Eb MOVXSX.Gv.Ew ]
is esc2; page = 1 & row = 11 & col = {2 to 7}
```

And the single opcode on row 12:

```
<patterns for integer opcodes>+= (<-U U->) [<-D>]
BSWAP is esc2; row = 12 & page = 1
```

## Operands and effective addresses

Intel operands and addresses are described in Section 25.2.1 and Figure 25-2 of the Pentium manual. Effective addresses use a ``Mod R/M" byte, the mod field of which determines the addressing mode. The Mod R/M byte also holds some bits that denote either a register operand or some extra parts of the opcode (as with the groupx instructions). Indexed addressing modes use an additional byte, called ``SIB," which holds a scale factor and index and base registers.

```
<field specs>+= (<-U U->) [<-D->]
fields of modrm (8) mod 6:7 reg_opcode 3:5 r_m 0:2
fields of sib (8) ss 6:7 index 3:5 base 0:2
```

```
<field specs>+= (<-U U->) [<-D->]
fieldinfo [ base index ] is
    [ names [ eAX eCX eDX eBX eSP eBP eSI eDI ] ]
fieldinfo ss is [ sparse [ "1" = 0, "2" = 1, "4" = 2, "8" = 3 ] ]
```

```
<placeholders>+= (<-U U->) [<-D->]
placeholder for modrm is HLT
placeholder for sib is HLT
```

We're faced with a specification problem because some Intel instructions accept only effective addresses that refer to operands in memory; register modes are not permitted. Most instructions, however, accept any kind of effective address. A good way to express this restriction would be with subtyping, but the toolkit doesn't implement subtyping, so instead we define two different constructor types: Mem to refer to effective addresses of operands in memory, and Eaddr to refer to any effective address. This separation requires the use of an identity constructor E to map Mem into Eaddr. Such a thing is bad enough in a specification, but these E's have to be used in application programs, too. The ugliness is justified because it confers protection against inadvertently using a register operand with an instruction that doesn't permit one.

```
<constructors for effective addresses>= (<-U U->)
relocatable d a
constructors
  Indir [reg] : Mem { reg != 4, reg != 5 } is mod = 0 & r_m = reg
  Disp8 d[reg] : Mem { reg != 4 } is mod = 1 & r_m = reg; i8 = d
  Disp32 d[reg] : Mem { reg != 4 } is mod = 2 & r_m = reg; i32 = d
  Abs32 a : Eaddr is mod = 0 & r_m = 5; i32 = a
  Reg reg : Eaddr is mod = 3 & r_m = reg
  Index [base][index * ss] : Mem { index != 4, base != 5 } is
    mod = 0 & r_m = 4; index & base & ss
  Index8 d[base][index * ss] : Mem { index != 4 } is
    mod = 1 & r_m = 4; index & base & ss; i8 = d
  Index32 d[base][index * ss] : Mem { index != 4 } is
    mod = 2 & r_m = 4; index & base & ss; i32 = d
  ShortIndex d[index * ss] : Mem { index != 4 } is
    mod = 0 & r_m = 4; index & base = 5 & ss; i32 = d
  E Mem : Eaddr is Mem
```

We'll eventually want to be able to keep only 32-bit constructors, to cut down on the time needed to generate encoding procedures.

```
<32-bit constructors>= (<-U) [D->]
Indir Disp32 Reg Index Index32 E
```

Now, this is good as far as it goes, but there are a couple of problems: no support for conditional assembly, and lots of constructors, which increases generation time. So let's get a bit clever:

```
<proposed constructors for effective addresses>=
constructors
  Indir0 [reg] { reg != 4, reg != 5 } is mod = 0 & r_m = reg
  Indir8 d[reg] { reg != 4 } is mod = 1 & r_m = reg; i8 = d
  Indir32 d[reg] { reg != 4 } is mod = 2 & r_m = reg; i32 = d
```

```

Index0  [base][index*ss] { index != 4, base != 5 } is
                        mod = 0 & r_m = 4; index & base    & ss
Index8  d[base][index*ss] { index != 4 } is
                        mod = 1 & r_m = 4; index & base    & ss; i8  = d
Index32 d[base][index*ss] { index != 4 } is
                        mod = 2 & r_m = 4; index & base    & ss; i32 = d

relocatable d
constructors
  Indir d[reg]          : Mem  when { d = 0 } is Indir0 (    reg)
                        when {      } is Indir8 (d, reg)
                        otherwise      is Indir32(d, reg)
  Index d[base][index*ss] : Mem  when { d = 0 } is Index0 (  base, index, ss)
                        when {      } is Index8 (d, base, index, ss)
                        otherwise      is Index32(d, base, index, ss)
  ShortIndex d[index*ss] : Mem { index != 4 } is
                        mod = 0 & r_m = 4; index & base = 5 & ss; i32 = d
  E  Mem : Eaddr is Mem
  Reg reg : Eaddr is mod = 3 & r_m = reg

discard Indir0 Indir8 Indir32 Index0 Index8 Index32

```

The only problem now is that it's no longer possible to split out the 32-bit constructors.

Immediate operands occupy whole tokens and have their own classes.

```

<field specs>+= (<-U U->) [<-D]
fields of I8  (8) i8  0:7
fields of I16 (16) i16 0:15
fields of I32 (32) i32 0:31

```

```

<placeholders>+= (<-U U->) [<-D]
placeholder for I8  is HLT
placeholder for I16 is HLT; HLT
placeholder for I32 is HLT; HLT; HLT; HLT

```

## Mod R/M opcodes

The Intel architecture offers the spectacle of putting some of the opcode bits in with the operands. The eight values of the `reg_opcode` field of the Mod R/M byte specify different opcodes, depending on the value of the opcode preceding the effective address. Most of these opcodes are notated by ``Groupx" in the manual. We use different sets of names for the values depending on what opcode precedes the effective-address specification. To make sure we don't mistakenly use a name like `INC.Eb` for `reg_opcode = 0` when the actual denotation is `INC.Ev`, we include in the specifications the opcode that must precede the Mod R/M byte---in this case, either `grp4` or `grp5`.

This kind of specification is conjoined below with specifications of opcodes and effective addresses. If

```
INC.Eb is grp4; reg_opcode = 0
```

then we might conjoin it with a `grp4` opcode followed by an effective address, writing:

```
INC.Eb Eaddr is (grp4; Eaddr) & INC.Eb
```

The conjunction of the explicit `grp4` with the `grp4` in the definition of `INC.Eb` ensures that `INC.Eb` is used correctly, since `grp4 & grp4 == grp4`. If we incorrectly used `grp5` when defining the `INC.Eb` constructor, `grp4 & grp5` would evaluate to a pattern that never matches anything, and the toolkit would complain.

We've glossed over an important detail in the definition of this constructor. Because conjunction distributes over concatenation, the right-hand side of the `INC.Eb` constructor winds up being equivalent to

```
grp4; (Eaddr & reg_opcode = 0).
```

This conjunction, unfortunately, breaks the rules of conjunction, which requires both patterns conjoined to have the same *shape*, or sequence of token classes. The conjunction is legal when `Eaddr` is an instance of `Indir` or `Reg`, because those effective addresses consist solely of one Mod R/M token, but the other modes contain more tokens, and their shapes don't match `reg_opcode = 0`. The solution is to relax the shape constraint by using the ellipsis operator. ```p ...`'' creates a pattern that is equivalent to `p`, except it is permissible to write `p ... & q` whenever `p`'s shape is a prefix of `q`'s shape. [ The ellipsis may also be used as a prefix operator on patterns, in which case `... p & q` is permissible



whenever p's shape is a suffix of q's shape. We haven't had occasion to use such patterns in machine descriptions, because most hardware decodes complex instructions from left to right. In the case at hand, every Eaddr begins with a Mod R/M token, so we can always write

```
Eaddr & reg_opcode = 0 ...
```

We've now covered enough detail to specify the Mod R/M or ``Groupx'' opcodes. Just to add some extra complexity, groups 1-3 include opcodes that denote different operand specifiers. For example, ADDi can denote an integer add of bytes (Eb.Ib), 16-byte or 32-byte words (Ev.Iv), or words and bytes (Ev.Ib), depending on the suffix added to the opcode. For each constructor created from the ADDi pattern, only one operand specifier is conjoined into the output pattern; the resulting pattern has just one non-contradictory disjunct, so the bits to be emitted are uniquely determined. The constructors for these patterns are defined in Section [\[->\]](#).

<pattern specs for other patterns>= (<-U U->) [D->]

patterns

```
arithI    is any of [ ADDi ORi ADCi SBBi ANDi SUBi XORi CMPi ],      # group 1
              which is (Eb.Ib | Ev.Iv | Ev.Ib); reg_opcode = {0 to 7} ...
bshifts   is B.Eb.1 | B.Eb.CL # D0 D2
vshifts   is B.Ev.1 | B.Ev.CL # D1 D3
immshifts is B.Eb.Ib | B.Ev.Ib # C0 C1
rot        is any of [ ROL ROR RCL RCR SHLSAL SHR SAR ],
              which is (bshifts | vshifts | immshifts);
              reg_opcode = {0 to 7} ...
grp3ops   is any of
[ TEST.Ib.Iv _ NOT NEG MUL.AL.eAX IMUL.AL.eAX DIV.AL.eAX IDIV.AL.eAX ],
              which is (grp3.Eb | grp3.Ev); reg_opcode = {0 to 7} ...
grp4ops   is any of [ INC.Eb DEC.Eb ],
              which is grp4; reg_opcode = [0 1] ...
grp5ops   is any of [ INC.Ev DEC.Ev CALL.Ep JMP.Ev JMP.Ep PUSH.Ev _ ],
              which is grp5; reg_opcode = {0 to 7} ...
grp6ops   is any of [ SLDT STR LLDT LTR VERR VERW _ _ ],
              which is grp6; reg_opcode = {0 to 7} ...
grp7ops   is any of [ SGDT SIDT LGDT LIDT SMSW _ LMSW INVLPG ],
              which is grp7; reg_opcode = {0 to 7} ...
bittestI  is any of [ BTi BTSi BTRi BTCi ],
              which is grp8; reg_opcode = {4 to 7} ...
CMPXCHG8B is          grp9; reg_opcode = 1 ...
```

## Operand-size and address-size prefixes

**[\*]** The Intel uses prefixes to distinguish 16- from 32-bit operands. The meaning of a prefix on the mode and on the setting of the *D* bit in the executable-segment descriptor (see pages 25-1ff of the Pentium manual). We assume here that *D*=1, making the default size 32 bits, but that assumption could be changed by reversing the definitions of *ow* and *od* given here. An application that wanted to be able to use both encodings would have to generate two sets of encoding procedures, perhaps using function pointers to switch back and forth.

In specifications, instructions with *b* suffixes (e.g., ``Eb,Gb'') use no prefix. Instructions with *v* suffixes (e.g., ``Ev,Gv'') begin with *ov*, which is followed by the rest of the instruction. When *ov* is used to build an opcode, this technique automatically creates two variants: *od*, a 32-bit variant with no prefix (epsilon) and *ow*, a 16-bit variant with prefix *OpPrefix*.

<prefix assignments>= (<-U U->) [D->]

```
patterns ow is OpPrefix
          od is epsilon
          ov is ow | od
```

The address prefix is similar, but we haven't figured out how it's to be used.

<prefix assignments>+= (<-U U->) [<-D]

```
patterns aw is AddrPrefix
          ad is epsilon
          av is aw | ad
```

## Floating-point opcodes

The specifications of floating-point opcodes consume many more tables, but it's not necessary to say much about them; the specification techniques needed are those we use above. We have defined patterns D9 through DF, which express



opcode values in hex; these are used in the specifications of the opcodes, so it seemed expedient to make them patterns, rather than continually writing something like opcode = 0xd8.

<pattern specs for other patterns>+= (<-U U->) [<-D->]

patterns

```
[ D8 D9 DA DB DC DD DE DF ] is ESC & col = {0 to 7}
[ FADD FMUL FCOM FCOMP FSUB FSUBR FDIV FDIVR ] is reg_opcode = {0 to 7}
[ FLD _ FST FSTP FLDENV FLDCW FSTENV FSTCW ] is reg_opcode = {0 to 7} ...
[ FNOP ] is D9; mod = 3 & reg_opcode = 2 & r_m = [0]
[ FCHS FABS _ FTST FXAM _ ] is D9; mod = 3 & reg_opcode = 4 & r_m = {0 to 7}
[ F2XM1 FYL2X FPTAN FPATAN FEXTRACT FPREM1 FDECSTP FINCSTP ]
is D9; mod = 3 & reg_opcode = 6 & r_m = {0 to 7}
FXCH is D9; mod = 3 & reg_opcode = 1
Fconstants is any of [ FLD1 FLDL2T FLDL2E FLDPI FLDLG2 FLDLN2 FLDZ _ ], which
is D9; mod = 3 & reg_opcode = 5 & r_m = {0 to 7}
[ FPREM FYL2XP1 FSQRT FSINCOS FRNDINT FSCALE FSIN FCOS ]
is D9; mod = 3 & reg_opcode = 7 & r_m = {0 to 7}
[ FIADD FIMUL FICOM FICOMP FISUB FISUBR FIDIV FIDIVR ] is reg_opcode = {0 to 7} ...
FUCOMPP is DA; mod = 3 & reg_opcode = 5 & r_m = 1
[ FILD _ FIST FISTP FBLD FLD.ext FBSTP FSTP.ext ] is reg_opcode = {0 to 7} ...
[ FCLEX FINIT ] is DB; mod = 3 & reg_opcode = 4 & r_m = [2 3]
[ FRSTOR _ FSAVE FSTSW ] is reg_opcode = {4 to 7} ...
[ FFREE _ FST.st FSTP.st FUCOM FUCOMP _ ] is mod = 3 & reg_opcode = {0 to 7}
[ FADDP _ FUBSRP FDIVRP FMULP _ FSUBP FDIVP ] is mod = 3 & reg_opcode = {0 to 7}
FCOMPP is DE; mod = 3 & reg_opcode = 3 & r_m = 1
FSTSW.AX is DF; mod = 3 & reg_opcode = 4 & r_m = 0
```

This next group of floating-point patterns define suffixes that we use on other opcodes, not actual opcodes.

<pattern specs for other patterns>+= (<-U U->) [<-D->]

patterns

```
.STi is DD; mod = 3
Fstack is any of [ .ST.STi .STi.St P.STi.ST ], which is [ D8 DC DE ]; mod = 3
Fint is any of [ .I32 .I16 ], which is [ DA DE ]
Fmem is any of [ .R32 .R64 ], which is [ D8 DC ]
FlsI is any of [ .lsI16 .lsI32 ], which is [ DF DB ]
FlsR is any of [ .lsR32 .lsR64 ], which is [ D9 DD ]
```

## Arithmetic instructions

[\*] There are eight arithmetic instructions which have many different modes and which are all treated alike. The regular modes are shown in the upper left corners of Figures <-> and <->; the immediate modes are the ``group 1" instructions (denoted here by arithI). This is the only part of the Intel specification we were able to factor very well, but it does give us 112 constructors in just a dozen lines, so it is worth doing.

<arithmetic constructors>+= (<-U U->)

constructors

```
arith^"iAL" i8! is arith & AL.Ib ; i8
arith^"iAX" i16! is ow; arith & eAX.Iv; i16
arith^"iEAX" i32! is od; arith & eAX.Iv; i32
arithI^"b" Eaddr, i8! is (Eb.Ib; Eaddr) & arithI; i8
arithI^"w" Eaddr, i16! is ow; (Ev.Iv; Eaddr) & arithI; i16
arithI^"d" Eaddr, i32! is od; (Ev.Iv; Eaddr) & arithI; i32
arithI^"ov^"b" Eaddr, i8! is ov; (Ev.Ib; Eaddr) & arithI; i8
arith^"mr^"b" Eaddr, reg8 is arith & Eb.Gb; Eaddr & reg_opcode = reg8 ...
arith^"mr^"ov Eaddr, reg is ov; arith & Ev.Gv; Eaddr & reg_opcode = reg ...
arith^"rmb" reg8, Eaddr is arith & Gb.Eb; Eaddr & reg_opcode = reg8 ...
arith^"rm^"ov reg, Eaddr is ov; arith & Gv.Ev; Eaddr & reg_opcode = reg ...
```

<32-bit constructors>+= (<-U) [<-D->]

```
arith^"iEAX" arithI^"d" arith^"mr^"od arith^"rm^"od
```

## Other instructions (in alphabetical order)

Trying to factor the non-arithmetic instructions proved a thankless task, so we've given almost all instructions merely in alphabetical order as they appear in Chapter 25 of the Pentium manual. There's a little bit of local factoring, as with

some bit operations.

It's not really appropriate to try to explain this part of the specification; the best way to read this section by comparing with the alphabetical section of the Pentium architecture manual [[cite intel:pentium](#)]. To generate a checker by the end of the millenium, we divide the spec into four parts and check each individually.

[<alphabetical constructors>=](#) ([<-U](#))

[<alphabetical constructors A-F>](#)

[<alphabetical constructors G-K>](#)

[<alphabetical constructors L-Q>](#)

[<alphabetical constructors R>](#)

[<alphabetical constructors S-Z>](#)

[<alphabetical constructors A-F>=](#) ([<-U](#) [U->](#)) [[D->](#)]

AAA

AAD is AAD; i8 = 10

AAM is AAM; i8 = 10

AAS

# ADC, ADD, AND are in arith group

ARPL Eaddr, reg16 is ARPL; Eaddr & reg\_opcode = reg16 ...

[<32-bit constructors>+=](#) ([<-U](#)) [[<-D->](#)]

AAA AAD AAM AAS

[<names of pentium constructors>=](#) ([U->](#)) [[D->](#)]

AAA AAD AAM AAS ARPL

Note that ARPL requires a 16-bit register, i.e. %ax, %bx, for its second operand.

The ``short'' variant of BOUND (boundw) requires a 16-bit register for its first operand.

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

constructors

BOUND<sup>ov</sup> reg, Mem is ov; BOUND; Mem & reg\_opcode = reg ...

BSF<sup>ov</sup> reg, Eaddr is ov; BSF; Eaddr & reg\_opcode = reg ...

BSR<sup>ov</sup> reg, Eaddr is ov; BSR; Eaddr & reg\_opcode = reg ...

BSWAP r32 is BSWAP & ... r32

BT<sup>ov</sup> Eaddr, reg is ov; BT; Eaddr & reg\_opcode = reg ...

BTi<sup>ov</sup> Eaddr, i8! is ov; (grp8; Eaddr) & BTi; i8

BTC<sup>ov</sup> Eaddr, reg is ov; BTC; Eaddr & reg\_opcode = reg ...

BTCi<sup>ov</sup> Eaddr, i8! is ov; (grp8; Eaddr) & BTCi; i8

BTR<sup>ov</sup> Eaddr, reg is ov; BTR; Eaddr & reg\_opcode = reg ...

BTRi<sup>ov</sup> Eaddr, i8! is ov; (grp8; Eaddr) & BTRi; i8

BTS<sup>ov</sup> Eaddr, reg is ov; BTS; Eaddr & reg\_opcode = reg ...

BTSi<sup>ov</sup> Eaddr, i8! is ov; (grp8; Eaddr) & BTSi; i8

[<names of pentium constructors>+=](#) ([U->](#)) [[<-D->](#)]

BOUND<sup>ov</sup> BSF<sup>ov</sup> BSR<sup>ov</sup> BSWAP BT<sup>ov</sup> BTi<sup>ov</sup> BTC<sup>ov</sup> BTCi<sup>ov</sup> BTR<sup>ov</sup> BTRi<sup>ov</sup> BTS<sup>ov</sup> BTSi<sup>ov</sup>

[<32-bit constructors>+=](#) ([<-U](#)) [[<-D->](#)]

BOUND<sup>od</sup> BSF<sup>od</sup> BSR<sup>od</sup> BSWAP BT<sup>od</sup> BTi<sup>od</sup> BTC<sup>od</sup> BTCi<sup>od</sup> BTR<sup>od</sup> BTRi<sup>od</sup> BTS<sup>od</sup> BTSi<sup>od</sup>

To deal with relative displacements, we set up constructors to compute them. The displacements are relative to the *end* of the word.

[<constructors for displacements>=](#) ([<-U](#) [U->](#))

constructors

rel8 reloc : Rel8 { reloc = L + i8! } is i8; L: epsilon

rel16 reloc : Rel16 { reloc = L + i16! } is i16; L: epsilon

rel32 reloc : Rel32 { reloc = L + i32! } is i32; L: epsilon

[<32-bit constructors>+=](#) ([<-U](#)) [[<-D->](#)]

rel32

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

CALL.Jv<sup>ow</sup> reloc is ow; CALL.Jv; rel16(reloc)

CALL.Jv<sup>od</sup> reloc is od; CALL.Jv; rel32(reloc)

CALL.Ep<sup>ov</sup> Mem is ov; (grp5; Mem) & CALL.Ep

```
CALL.aP^ow CS:" IP is ow; CALL.aP; i16 = CS; i16 = IP
CALL.aP^od CS:" IP is od; CALL.aP; i16 = CS; i32 = IP
CALL.Ev^ov Mem is ov; (grp5; Mem) & CALL.Ev
CBW is ow; CBW
CWDE is od; CBW
CLC
CLD
CLI
CLTS
CMC
```

The Linux assembler doesn't support multiple segments, so the CALL opcodes that take a code segment and offset are discarded when generating assembly code for a Linux assembler.

[<pentium-linux.spec>= \(U-> U->\) \[D->\]](#)  
discard CALL.aP^ow CALL.aP^od

[<alphabetical constructors A-F>+= \(<-U U->\) \[<-D->\]](#)  
# CMP is in the arith group  
CMPSB^av is av; CMPSB  
CMPSv^ov^av is (av; ov | ov; av); CMPSv  
CMPXCHG.Eb.Gb Eaddr, reg is CMPXCHG.Eb.Gb; Eaddr & reg\_opcode = reg ...  
CMPXCHG.Ev.Gv^ov Eaddr, reg is ov; CMPXCHG.Ev.Gv; Eaddr & reg\_opcode = reg ...  
CMPXCHG8B Mem is (grp9; Mem) & CMPXCHG8B  
CPUID  
CWD is ow; CWDQ  
CDQ is od; CWDQ

[<32-bit constructors>+= \(<-U\) \[<-D->\]](#)  
CALL.Jv^od CALL.Ep^od CALL.Ev^od CALL.aP^od CBW CWDE CLC CLD CLI CLTS CMC  
CMPSv^od^av CMPXCHG.Ev.Gv^od CWD CDQ

CMPXCHG8B and CPUID are Pentium instructions so we can't check them on a 486.

[<names of pentium constructors>+= \(U->\) \[<-D->\]](#)  
CMPXCHG8B CPUID CWD

[<alphabetical constructors A-F>+= \(<-U U->\) \[<-D->\]](#)  
DAA  
DAS  
DEC.Eb Eaddr is (grp4; Eaddr) & DEC.Eb  
DEC.Ev^ov Eaddr is ov; (grp5; Eaddr) & DEC.Ev  
DEC^ov r32 is ov; DEC & r32  
DIV^"AL" Eaddr is (grp3.Eb; Eaddr) & DIV.AL.eAX  
DIV^"AX" Eaddr is ow; (grp3.Ev; Eaddr) & DIV.AL.eAX  
DIV^"eAX" Eaddr is od; (grp3.Ev; Eaddr) & DIV.AL.eAX

[<32-bit constructors>+= \(<-U\) \[<-D->\]](#)  
DAA DAS DEC.Ev^od DEC^od DIV^"eAX"

[<names of pentium constructors>+= \(U->\) \[<-D->\]](#)  
DAA DAS DEC.Eb

[<alphabetical constructors A-F>+= \(<-U U->\) \[<-D->\]](#)  
ENTER i16, i8! is ENTER; i16; i8  
F2XM1

[<32-bit constructors>+= \(<-U\) \[<-D->\]](#)  
ENTER F2XM1

[<names of pentium constructors>+= \(U->\) \[<-D->\]](#)  
ENTER F2XM1

The first use of the left-hand side ellipsis is used in the definition of FADD^Fstack. The left ellipsis denotes that the pattern (FADD & r\_m = idx) must be a legal suffix of the pattern Fstack.

[<alphabetical constructors A-F>+= \(<-U U->\) \[<-D->\]](#)  
FABS

FADD^Fmem Mem is Fmem; Mem & FADD ...  
FADD^Fstack idx is Fstack & ... (FADD & r\_m = idx)

[<floating-point constructors>+=](#) ([<-U](#) [<-U](#) [U->](#)) [[<-D->](#)]  
FABS FADD^Fmem FADD^Fstack

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]  
FIADD^Fint Mem is Fint; Mem & FIADD ...  
FBLD Mem is DF; Mem & FBLD  
FBSTP Mem is DF; Mem & FBSTP

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]  
FCHS

[<floating-point constructors>+=](#) ([<-U](#) [<-U](#) [U->](#)) [[<-D->](#)]  
FCHS

[<synthetics with WAIT>=](#) ([U->](#)) [[<-D->](#)]  
FCLEX is WAIT; FCLEX

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]  
FNCLEX is FCLEX

[<pattern specs for other patterns>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]  
patterns FCOMs is FCOM | FCOMP

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]  
constructors  
FCOMs^Fmem Mem is Fmem; Mem & FCOMs ... # includes FICOM, FICOMP  
FCOMs^.ST.STi idx is .ST.STi & ... (FCOMs & r\_m = idx)  
FCOMPP  
FCOS

[<floating-point constructors>+=](#) ([<-U](#) [<-U](#) [U->](#)) [[<-D->](#)]  
FCOMs^Fmem FCOMs^.ST.STi FCOMPP FCOS

[<pentium-linux.spec>+=](#) ([U->](#) [U->](#)) [[<-D->](#)]  
discard FCOMs^.ST.STi

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]  
FDECSTP  
FDIV^Fmem Mem is Fmem; Mem & FDIV ...  
FDIV^Fstack idx is Fstack & ... (FDIV & r\_m = idx)  
FDIVR^Fmem Mem is Fmem; Mem & FDIVR ...  
FDIVR^Fstack idx is Fstack & ... (FDIVR & r\_m = idx)

[<floating-point constructors>+=](#) ([<-U](#) [<-U](#) [U->](#)) [[<-D->](#)]  
FDECSTP FDIV^Fmem FDIV^Fstack FDIVR^Fmem FDIVR^Fstack

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]  
FFREE idx is DD; FFREE & r\_m = idx

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]  
FICOM^Fint Mem is Fint; Mem & FICOM  
FICOMP^Fint Mem is Fint; Mem & FICOMP  
FILD^FlsI Mem is FlsI; Mem & FILD  
FILD64 Mem is DF; Mem & FLD.ext ...  
FINIT

[<pattern specs for other patterns>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]  
patterns FISTs is FIST | FISTP

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]  
constructors  
FISTs^FlsI Mem is FlsI; Mem & FISTs  
FISTP64 Mem is DF; Mem & FSTP.ext

[<floating-point constructors>+=](#) ([<-U](#) [<-U](#) [U->](#)) [[<-D->](#)]  
FICOM^Fint FICOMP^Fint FICOM16 FICOM32 FICOMP16 FICOMP32

FILD<sup>FlsI</sup> FILD64 FINIT FISTs<sup>FlsI</sup> FISTP64

*EDIT ME? There is no obvious way to reference indirectly the stack pointer, i.e., (%esp), because the encoding of this addressing mode is an escape that indicates an sib byte follows the mod-rm byte. The only way to construct this address is to build a sib byte that ignores its index field. Page 26-7 specifies how to do this. Luckily, we only need (%esp) for a few floating-point instructions that use the address on the top of the stack and then pops it off. Although this addressing mode is technically an Eaddr, we specify it individually because virutally no other instruction uses it.*

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

FLD<sup>FlsR</sup> Mem is FlsR; Mem & FLD  
 FLD80 Mem is DB; Mem & FLD.ext ...  
 FLD.STi idx is D9; mod = 3 & FLD & r\_m = idx  
 Fconstants  
 FLDCW Mem is D9; Mem & FLDCW  
 FLDENV Mem is D9; Mem & FLDENV

[<floating-point constructors>+=](#) ([<-U](#) [<-U](#) [U->](#)) [[<-D->](#)]

FLD<sup>FlsR</sup> FLD80 FLD.STi Fconstants FLDCW FLDENV

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

FMUL<sup>Fmem</sup> Mem is Fmem; Mem & FMUL ...  
 FMUL<sup>Fstack</sup> idx is Fstack & ... (FMUL & r\_m = idx)

[<floating-point constructors>+=](#) ([<-U](#) [<-U](#) [U->](#)) [[<-D->](#)]

FMUL<sup>Fmem</sup> FMUL<sup>Fstack</sup>

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

FNOP

[<floating-point constructors>+=](#) ([<-U](#) [<-U](#) [U->](#)) [[<-D->](#)]

FNOP

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

FPATAN  
 FPREM  
 FPREM1  
 FPTAN

[<floating-point constructors>+=](#) ([<-U](#) [<-U](#) [U->](#)) [[<-D->](#)]

FPATAN FPREM FPREM1 FPTAN

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

FRNDINT  
 FRSTOR Mem is DD; Mem & FRSTOR

[<floating-point constructors>+=](#) ([<-U](#) [<-U](#) [U->](#)) [[<-D->](#)]

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

FNSAVE Mem is DD; Mem & FSAVE

[<floating-point constructors>+=](#) ([<-U](#) [<-U](#) [U->](#)) [[<-D->](#)]

[<synthetics with WAIT>+=](#) ([U->](#)) [[<-D->](#)]

FSAVE Mem is WAIT; FNSAVE(Mem)

[<floating-point constructors>+=](#) ([<-U](#) [<-U](#) [U->](#)) [[<-D->](#)]

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

FSCALE  
 FSIN  
 FSINCOS  
 FSQRT

[<pattern specs for other patterns>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

patterns  
 FSTs is FST | FSTP  
 FSTs.st is FST.st | FSTP.st

[<alphabetical constructors A-F>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

```

constructors
FSTs^FlsR      Mem is FlsR; Mem & FSTs
FSTP80         Mem is DB; Mem & FSTP.ext
FSTs.st^.STi   idx is .STi & ... (FSTs.st & r_m = idx)
FSTCW          Mem is D9; Mem & FSTCW

<synthetics with WAIT>+= (U->) [<-D->]
FNSTCW         Mem is WAIT; FSTCW(Mem)

<floating-point constructors>+= (<-U <-U U->) [<-D->]
FSCALE FSIN FSINCOS FSQRT FSTs^FlsR FSTP80 FSTs.st^.STi FSTCW FNSTCW

<alphabetical constructors A-F>+= (<-U U->) [<-D->]
FSTENV         Mem is D9; Mem & FSTENV

<synthetics with WAIT>+= (U->) [<-D->]
FNSTENV        Mem is WAIT; FSTENV(Mem)

<floating-point constructors>+= (<-U <-U U->) [<-D->]
FSTENV FNSTENV

<alphabetical constructors A-F>+= (<-U U->) [<-D->]
FSTSW          Mem is DD; Mem & FSTSW
FSTSW.AX

<synthetics with WAIT>+= (U->) [<-D->]
FNSTSW         Mem is WAIT; FSTSW(Mem)
FNSTSW.AX      is WAIT; FSTSW.AX()

<floating-point constructors>+= (<-U <-U U->) [<-D->]
FSTSW FSTSW.AX FNSTSW FNSTSW.AX

<alphabetical constructors A-F>+= (<-U U->) [<-D->]
FSUB^Fmem      Mem is Fmem; Mem & FSUB ...
FSUB^Fstack    idx is Fstack & ... (FSUB & r_m = idx)
FSUBR^Fmem     Mem is Fmem; Mem & FSUBR ...
FSUBR^Fstack   idx is Fstack & ... (FSUBR & r_m = idx)

<floating-point constructors>+= (<-U <-U U->) [<-D->]
FSUB^Fmem FSUB^Fstack FSUBR^Fmem FSUBR^Fstack

<alphabetical constructors A-F>+= (<-U U->) [<-D->]
FTST

<floating-point constructors>+= (<-U <-U U->) [<-D->]
FTST

<pattern specs for other patterns>+= (<-U U->) [<-D->]
patterns FUCOMs is FUCOM | FUCOMP

<alphabetical constructors A-F>+= (<-U U->) [<-D->]
constructors
FUCOMs idx is DD; FUCOMs & r_m = idx
FUCOMPP

<floating-point constructors>+= (<-U <-U U->) [<-D->]
FUCOMs FUCOMPP

<synthetics with WAIT>+= (U->) [<-D]
FWAIT is WAIT

<floating-point constructors>+= (<-U <-U U->) [<-D->]
FWAIT

<alphabetical constructors A-F>+= (<-U U->) [<-D->]
FXAM
FXCH idx is FXCH & ... r_m = idx
EXTRACT

```

<floating-point constructors>+= (<-U <-U U->) [<-D->]

FXAM FXCH FXTRACT

<alphabetical constructors A-F>+= (<-U U->) [<-D]

FYL2X  
FYL2XP1

<floating-point constructors>+= (<-U <-U U->) [<-D]

FYL2X FYL2XP1

<names of pentium constructors>+= (U->) [<-D]

FIADD^Fint FBLD FBSTP FNCLEX FCOS FDECSTP FDIVR^Fmem FIDIV^Fint FIDIVR^Fint  
FICOM^Fint FICOMP^Fint FILD64 FINCSTP FINIT FISTP64 FLD80 Fconstants  
FLDCW FLDENV FMUL^Fint FNOP FPATAN FPREM FPREM1 FPTAN FRNDINT  
FRSTOR FSCALE FSIN FSINCOS FSQRT FSUBR^Fmem FISUB^Fint FISUBR^Fint  
FTST FUCOMs FUCOMPP FXAM FXCH FXTRACT FYL2X FYL2XP1

<alphabetical constructors G-K>= (<-U U->) [D->]

HLT

<32-bit constructors>+= (<-U) [<-D->]

HLT

<alphabetical constructors G-K>+= (<-U U->) [<-D->]

IDIV Eaddr is (grp3.Eb; Eaddr) & IDIV.AL.eAX  
IDIV^"AX" Eaddr is ow; (grp3.Ev; Eaddr) & IDIV.AL.eAX  
IDIV^"eAX" Eaddr is od; (grp3.Ev; Eaddr) & IDIV.AL.eAX  
IMULb Eaddr is (grp3.Eb; Eaddr) & IMUL.AL.eAX  
IMUL^ov Eaddr is ov; (grp3.Ev; Eaddr) & IMUL.AL.eAX  
IMULrm^ov reg, Eaddr is ov; IMUL.Gv.Ev; Eaddr & reg\_opcode = reg ...  
IMUL.Ib^ov reg, Eaddr, i8! is ov; IMUL.Ib; Eaddr & reg\_opcode = reg ... ; i8  
IMUL.Iv^"w" reg, Eaddr, i16! is ow; IMUL.Iv; Eaddr & reg\_opcode = reg ... ; i16  
IMUL.Iv^"d" reg, Eaddr, i32! is od; IMUL.Iv; Eaddr & reg\_opcode = reg ... ; i32  
IN.AL.Ib i8! is IN.AL.Ib; i8  
IN.eAX.Ib^ov i8! is ov; IN.eAX.Ib; i8  
IN.AL.DX  
IN.eAX.DX^ov is ov; IN.eAX.DX  
INC.Eb Eaddr is (grp4; Eaddr) & INC.Eb  
INC.Ev^ov Eaddr is ov; (grp5; Eaddr) & INC.Ev  
INC^ov r32 is ov; INC & r32  
INSB  
INSv^ov is ov; INSv  
INT3  
INT.Ib i8! is INT.Ib; i8  
INT0  
INVD  
INVLPG Mem is (grp7; Mem) & INVLPG  
IRET

<32-bit constructors>+= (<-U) [<-D->]

IDIV^"eAX" IMUL^od IMULrm^od IMUL.Iv^"d" IN.eAX.DX^od INC.Ev^od INC^od  
INSv^od INT3 INT.Ib INTO INVD INVLPG IRET

<alphabetical constructors G-K>+= (<-U U->) [<-D->]

Jb^cond reloc is Jb & cond; rel8(reloc)  
Jv^cond^ow reloc is ow; Jv & ... cond; rel16(reloc)  
Jv^cond^od reloc is od; Jv & ... cond; rel32(reloc)  
JMP.Jb reloc is JMP.Jb; rel8(reloc)  
JMP.Jv^ow reloc is ow; JMP.Jv; rel16(reloc)  
JMP.Jv^od reloc is od; JMP.Jv; rel32(reloc)  
JMP.Ap^ow CS, IP is ow; JMP.Ap; i16 = CS; i16 = IP  
JMP.Ap^od CS, IP is od; JMP.Ap; i16 = CS; i32 = IP  
JMP.Ev^ov Eaddr is ov; (grp5; Eaddr) & JMP.Ev  
JMP.Ep^ov Mem is ov; (grp5; Mem) & JMP.Ep

<32-bit constructors>+= (<-U) [<-D->]

Jv^cond^od JMP.Jv^od JMP.Ap^od JMP.Ev^od JMP.Ep^od



[<pentium-linux.spec>+= \(U-> U->\) \[<-D->\]](#)

discard JMP.Ap<sup>ow</sup> JMP.Ap<sup>ov</sup>

[<alphabetical constructors L-Q>+= \(<-U U->\) \[D->\]](#)

LAHF

LAR<sup>ov</sup> reg, Eaddr is ov; LAR; Eaddr & reg\_opcode = reg ...

[<pattern specs for other patterns>+= \(<-U U->\) \[<-D->\]](#)

patterns lfp is LDS | LES | LFS | LGS | LSS

[<alphabetical constructors G-K>+= \(<-U U->\) \[<-D>\]](#)

constructors

lfp<sup>ov</sup> reg, Mem is ov; lfp; Mem & reg\_opcode = reg ...

LEA<sup>ov</sup> reg, Mem is ov; LEA; Mem & reg\_opcode = reg ...

LEAVE

LGDT Mem is (grp7; Mem) & LGDT

LIDT Mem is (grp7; Mem) & LIDT

LLDT Eaddr is (grp6; Eaddr) & LLDT

LMSW Eaddr is (grp7; Eaddr) & LMSW

LOCK

LODSB

LODSv<sup>ov</sup> is ov; LODSv

[<pattern specs for other patterns>+= \(<-U U->\) \[<-D->\]](#)

patterns LOOPs is LOOP | LOOPE | LOOPNE

[<alphabetical constructors L-Q>+= \(<-U U->\) \[<-D->\]](#)

constructors

LOOPs<sup>ov</sup> reloc is ov; LOOPs; rel8(reloc)

LSL<sup>ov</sup> reg, Eaddr is ov; LSL; Eaddr & reg\_opcode = reg ...

LTR Eaddr is (grp6; Eaddr) & LTR

[<32-bit constructors>+= \(<-U\) \[<-D->\]](#)

LAHF LAR<sup>ov</sup> lfp<sup>ov</sup> LEA<sup>ov</sup> LEAVE LGDT LIDT LLDT LMSW LOCK LODSB LODSv<sup>ov</sup>

[<alphabetical constructors L-Q>+= \(<-U U->\) \[<-D->\]](#)

MOV<sup>^</sup>"mrb" Eaddr, reg is MOV & Eb.Gb; Eaddr & reg\_opcode = reg ...

MOV<sup>^</sup>"mr"<sup>ov</sup> Eaddr, reg is ov; MOV & Ev.Gv; Eaddr & reg\_opcode = reg ...

MOV<sup>^</sup>"rmb" reg, Eaddr is MOV & Gb.Eb; Eaddr & reg\_opcode = reg ...

MOV<sup>^</sup>"rm"<sup>ov</sup> reg, Eaddr is ov; MOV & Gv.Ev; Eaddr & reg\_opcode = reg ...

MOV.Ew.Sw Mem, srl6 is ov; MOV.Ew.Sw; Mem & reg\_opcode = srl6 ...

MOV.Sw.Ew Mem, srl6 is MOV.Sw.Ew; Mem & reg\_opcode = srl6 ...

# assume 32-bit address mode

MOV.AL.Ob offset is MOV.AL.Ob; i32 = offset

MOV.eAX.0v<sup>ov</sup> offset is ov; MOV.eAX.0v; i32 = offset

MOV.Ob.AL offset is MOV.Ob.AL; i32 = offset

MOV.0v.eAX<sup>ov</sup> offset is ov; MOV.0v.eAX; i32 = offset

MOVib r8, i8! is MOVib & r8; i8

MOViv r32, i16! is MOViv & r32; i16

MOVid r32, i32! is od; MOViv & r32; i32

MOV.Eb.Ib Eaddr, i8! is MOV.Eb.Ib; Eaddr & reg\_opcode = 0 ...; i8

MOV.Ev.Iv<sup>ow</sup> Eaddr, i16! is ov; MOV.Ev.Iv; Eaddr & reg\_opcode = 0 ...; i16

MOV.Ev.Iv<sup>od</sup> Eaddr, i32! is od; MOV.Ev.Iv; Eaddr & reg\_opcode = 0 ...; i32

MOV.Cd.Rd cr, reg is MOV.Cd.Rd; mod = 3 & r\_m = reg & reg\_opcode = cr

MOV.Rd.Cd reg, cr is MOV.Rd.Cd; mod = 3 & r\_m = reg & reg\_opcode = cr

MOV.Dd.Rd dr, reg is MOV.Dd.Rd; mod = 3 & r\_m = reg & reg\_opcode = dr

MOV.Rd.Dd reg, dr is MOV.Rd.Dd; mod = 3 & r\_m = reg & reg\_opcode = dr

MOVSB

MOVSv<sup>ov</sup> is ov; MOVSv

MOV SX.Gv.Eb<sup>ov</sup> r32, Eaddr is ov; MOV SX.Gv.Eb; Eaddr & reg\_opcode = r32 ...

MOV SX.Gv.Ew r16, Eaddr is MOV SX.Gv.Ew; Eaddr & reg\_opcode = r16 ...

MOVZX.Gv.Eb<sup>ov</sup> r32, Eaddr is ov; MOVZX.Gv.Eb; Eaddr & reg\_opcode = r32 ...

MOVZX.Gv.Ew r16, Eaddr is MOVZX.Gv.Ew; Eaddr & reg\_opcode = r16 ...

MUL.AL Eaddr is (grp3.Eb; Eaddr) & MUL.AL.eAX

MUL.AX<sup>ov</sup> Eaddr is ov; (grp3.Ev; Eaddr) & MUL.AL.eAX

[<32-bit constructors>+= \(<-U\) \[<-D->\]](#)

```
MOV^"mrb" MOV^"mr"^ov MOV^"rmb" MOV^"rm"^ov
MOV.Ew.Sw MOV.Sw.Ew MOV.AL.Ob MOV.eAX.Ov^ov MOV.Ob.AL
MOV.Ov.eAX^ov MOVib MOViw MOVid MOV.Eb.Ib MOV.Ev.Iv^ow
MOV.Ev.Iv^od MOVSB MOVSV^ov MOVVSX.Gv.Eb^od
MOVVSX.Gv.Ew MOVZX.Gv.Eb^od MOVZX.Gv.Ew MUL.AL MUL.AX^ov
```

<pentium-linux.spec>+= (U-> U->) [<-D->]

```
discard MOVVSX.Gv.Ew MOVZX.Gv.Eb^ov MOVZX.Gv.Ew MOV.Ew.Sw
```

<alphabetical constructors L-Q>+= (<-U U->) [<-D->]

```
NEGb Eaddr is (grp3.Eb; Eaddr) & NEG
NEG^ov Eaddr is ov; (grp3.Ev; Eaddr) & NEG
NOP
NOTb Eaddr is (grp3.Eb; Eaddr) & NOT
NOT^ov Eaddr is ov; (grp3.Ev; Eaddr) & NOT
```

<32-bit constructors>+= (<-U) [<-D->]

```
NEGb NEG^ov NOP NOTb NOT^ov
```

<alphabetical constructors L-Q>+= (<-U U->) [<-D->]

```
# OR is in the arith group
OUT.Ib.AL i8! is OUT.Ib.AL; i8
OUT.Ib.eAX^ov i8! is ov; OUT.Ib.eAX; i8
OUT.DX.AL
OUT.DX.eAX^ov is ov; OUT.DX.eAX
OUTSB
OUTSV^ov is ov; OUTSV
```

<32-bit constructors>+= (<-U) [<-D->]

```
OUT.Ib.AL OUT.Ib.eAX^ov OUT.DX.AL OUT.DX.eAX^ov OUTSB OUTSV^ov
```

<alphabetical constructors L-Q>+= (<-U U->) [<-D->]

```
POP.Ev^ov Mem is ov; POP.Ev; Mem & reg_opcode = 0 ...
POP^ov r32 is ov; POP & r32
```

<pattern specs for other patterns>+= (<-U U->) [<-D->]

```
patterns POPs is POP.ES | POP.SS | POP.DS | POP.FS | POP.GS
POPv is POPA | POPF
```

<alphabetical constructors L-Q>+= (<-U U->) [<-D->]

```
constructors
POPs
POPv^ov is ov; POPv
PUSH.Ev^ov Eaddr is ov; (grp5; Eaddr) & PUSH.Ev
PUSH^ov r32 is ov; PUSH & r32
PUSH.Ib i8! is PUSH.Ib; i8
PUSH.Iv^ow i16! is ov; PUSH.Iv; i16
PUSH.Iv^od i32! is od; PUSH.Iv; i32
```

<pattern specs for other patterns>+= (<-U U->) [<-D->]

```
patterns PUSHs is PUSH.CS | PUSH.SS | PUSH.DS | PUSH.ES | PUSH.FS | PUSH.GS
PUSHv is PUSHA | PUSHF
```

<alphabetical constructors L-Q>+= (<-U U->) [<-D->]

```
constructors
PUSHs
PUSHv^ov is ov; PUSHv
```

<32-bit constructors>+= (<-U) [<-D->]

```
POPs POPv^ov PUSH.Ev^ov PUSH^ov PUSH.Ib PUSH.Iv^ow PUSH.Iv^od PUSHs PUSHv^ov
```

<alphabetical constructors R>+= (<-U U->) [D->]

```
# ROL ROR RCL RCR SHLSAL SHR SAR
rot^bshifts Eaddr is (bshifts; Eaddr) & rot
rot^vshifts^ov Eaddr is ov; (vshifts; Eaddr) & rot
rot^B.Eb.Ib Eaddr, i8! is (B.Eb.Ib; Eaddr) & rot; i8
rot^B.Ev.Ib^ov Eaddr, i8! is ov; (B.Ev.Ib; Eaddr) & rot; i8
```

<32-bit constructors>+= (<-U) [<-D->]  
rot^bshifts rot^vshifts^ov rot^B.Ev.Ib^ov

<alphabetical constructors R>+= (<-U U->) [<-D]

RDMSR  
REP  
REPNE  
RET  
RET.far  
RET.Iw i16 is RET.Iw; i16  
RET.far.Iw i16 is RET.far.Iw; i16  
RSM

<32-bit constructors>+= (<-U) [<-D->]  
RDMSR REP REPNE RET RET.far RET.Iw RET.far.Iw RSM

<pentium-linux.spec>+= (U-> U->) [<-D->]  
discard RDMSR

<alphabetical constructors S-Z>+= (<-U U->) [D->]

SAHF  
# SAL SAR SHL SR above with rot  
# SBB is in the arith group  
SCASB  
SCASv^ov is ov; SCASv  
## SETb^cond Mem is SETb & ... cond; Mem  
SETb^cond Eaddr is SETb & ... cond; Eaddr  
SGDT Mem is (grp7; Mem) & SGDT  
SIDT Mem is (grp7; Mem) & SIDT

<32-bit constructors>+= (<-U) [<-D->]  
SCASB SCASv^ov SETb^cond SGDT SIDT

<pattern specs for other patterns>+= (<-U U->) [<-D]

patterns shdIb is SHRD.Ib | SHLD.Ib  
shdCL is SHRD.CL | SHLD.CL

<alphabetical constructors S-Z>+= (<-U U->) [<-D->]

constructors  
shdIb^ov Eaddr, reg, count is ov; shdIb; Eaddr & reg\_opcode = reg ... ; i8 = count  
shdCL^ov Eaddr, reg, "CL" is ov; shdCL; Eaddr & reg\_opcode = reg ...  
SLDT Eaddr is (grp6; Eaddr) & SLDT  
SMSW Eaddr is (grp7; Eaddr) & SMSW  
STC  
STD  
STI  
STOSB  
STOSv^ov is ov; STOSv  
STR Mem is (grp6; Mem) & STR  
# SUB is in the arith group

<32-bit constructors>+= (<-U) [<-D->]  
shdIb^ov shdCL^ov SLDT SMSW STC STD STI STOSB STOSv^ov STR

<alphabetical constructors S-Z>+= (<-U U->) [<-D->]

TEST.AL.Ib i8 is TEST.AL.Ib; i8  
TEST.eAX.Iv^ow i16 is ow; TEST.eAX.Iv; i16  
TEST.eAX.Iv^od i32 is od; TEST.eAX.Iv; i32  
TEST.Eb.Ib Eaddr, i8 is (grp3.Eb; Eaddr) & TEST.Ib.Iv; i8  
TEST.Ew.Iw Eaddr, i16 is ow; (grp3.Ev; Eaddr) & TEST.Ib.Iv; i16  
TEST.Ed.Id Eaddr, i32 is od; (grp3.Ev; Eaddr) & TEST.Ib.Iv; i32  
TEST.Eb.Gb Eaddr, reg is TEST.Eb.Gb; Eaddr & reg\_opcode = reg ...  
TEST.Ev.Gv^ov Eaddr, reg is ov; TEST.Ev.Gv; Eaddr & reg\_opcode = reg ...

<32-bit constructors>+= (<-U) [<-D->]

TEST.AL.Ib TEST.eAX.Iv^ow TEST.eAX.Iv^od TEST.Eb.Ib TEST.Ew.Iw  
TEST.Ed.Id TEST.Eb.Gb TEST.Ev.Gv^ov

[<alphabetical constructors S-Z>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

VERR Eaddr is (grp6; Eaddr) & VERR  
VERW Eaddr is (grp6; Eaddr) & VERW

[<32-bit constructors>+=](#) ([<-U](#)) [[<-D->](#)]

VERR VERW

[<alphabetical constructors S-Z>+=](#) ([<-U](#) [U->](#)) [[<-D->](#)]

WAIT  
WBINVD  
WRMSR

[<32-bit constructors>+=](#) ([<-U](#)) [[<-D->](#)]

WAIT WBINVD WRMSR

[<pentium-linux.spec>+=](#) ([U->](#) [U->](#)) [[<-D](#)]

discard WRMSR

[<alphabetical constructors S-Z>+=](#) ([<-U](#) [U->](#)) [[<-D](#)]

XADD.Eb.Gb Eaddr, reg is XADD.Eb.Gb; Eaddr & reg\_opcode = reg ...  
XADD.Ev.Gv^ov Eaddr, reg is ov; XADD.Ev.Gv; Eaddr & reg\_opcode = reg ...  
XCHG^"eAX"^ov r32 is ov; XCHG & r32  
XCHG.Eb.Gb Eaddr, reg is XCHG.Eb.Gb; Eaddr & reg\_opcode = reg ...  
XCHG.Ev.Gv^ov Eaddr, reg is ov; XCHG.Ev.Gv; Eaddr & reg\_opcode = reg ...  
XLATB is XLAT  
# XOR is in the arith group

[<32-bit constructors>+=](#) ([<-U](#)) [[<-D](#)]

XADD.Eb.Gb XADD.Ev.Gv^ov XCHG^"eAX"^ov XCHG.Eb.Gb XCHG.Ev.Gv^ov XLATB

## Synthetic instructions

The only synthetics we've identified are those that are preceded by a WAIT instruction, which we've included in the alphabetical list above, since they're all given together in Chapter 25.

The Gnu-Linux assembler does not recognize instructions prefixed by WAIT as synthetics. For example, wait; fclex disassembles as fclex; fclex disassembles as fnclex.

[<pentium-synth.spec>=](#)

constructors

[<synthetics with WAIT>](#)

mld needs some synthetics to help it deal with overloaded operators. This is a hack but better than writing them all out by hand.

[<mld-pentium.spec>=](#) [[D->](#)]

constructors  
\_call\_l reloc : Function is call\_jvod(reloc)  
\_call\_rm Mem : Function is call\_evod(Mem)  
callfn Function is Function

The following constructor is used to emit relocatable addresses. It is the same as that defined for the MIPS and SPARC.

[<mld-pentium.spec>+=](#) [[<-D](#)]

fields of addrtoken (32) addr32 0:31  
placeholder for addrtoken is addr32 = 7  
constructors  
emit\_raddr reloc is addr32 = reloc

## Assembly Specification for Gnu-Linux Assembler

Some assemblers overload instruction names, using context to determine which instruction is meant. For example, the Pentium add can mean any of five different instructions, depending on the sizes and locations of the operands. The toolkit cannot do this kind of overloading; it must use different names for different instructions (constructors). The reason is that the toolkit must generate a different encoding procedure for each instruction, and in most programming languages, different procedures must have different names. Even in languages that do support overloading, we might not be able to use the same name, because the name-resolution mechanisms used in programming languages are typically type-based and quite different from what an assembler uses (LR parsing).

We solve this problem by requiring each constructor to have a different name. Typically, the specification writer distinguishes variants by adding suffixes to the base name of the constructor. For example, the Pentium add instructions include constructors called addb, addib, addiowb, addmrb, and addrmb. To get from these names back to assembly language, we have to define an appropriate mapping. These mappings are defined separately from the main specification, because different vendors use different syntaxes for their assembly languages.

An assembly specification includes three parts: a mapping of constructor names to assembly opcodes, the assembly format for each constructor operand, and the assembly syntax for each constructor. We use the specification for the assembly language supported by the Gnu-Linux assembler to illustrate assembly specifications. First, we describe assembly name mappings, then operand format and constructor assembly syntax.

### Assembly names for opcodes

A constructor's name may contain multiple parts derived from pattern names and constant strings. Each part may or may not contribute to the assembly name. For example, the constructor name add<sup>^</sup>"mrb" contains two parts, the first derived from the pattern add and the second from the string "mrb". The assembly name for this constructor is addb, so the pattern add contributes its name and the suffix "mrb" is mapped to the string "b". This example illustrates how the constructor addmrb has a more specific name to disambiguate it from other overloaded variants of addb.

assembly opcode introduces mappings from complete constructor names (strings) to their assembly names (strings). assembly component introduces mappings from parts of constructor names to their assembly names. We provide a component-wise mapping, because it improves factoring of assembly names among constructors that share common suffixes and prefixes. For example, the suffix B.Eb.Ib always maps to b in every constructor name where it appears.

There is redundancy in the mapping of constructor names, so we use a regular expression syntax to group related names. The regular expression syntax is the same as the syntax for C-shell "globbing" expressions [cite joy:c-shell]. If a complete name mapping exists for a constructor name, it is applied first. If no complete mapping exists, mappings are applied individually to *each* part of a constructor's name and the resulting strings are concatenated into the complete assembly name. For example, the mappings applied to the parts of the constructor name add<sup>^</sup>"mrb" are:

#### <example mappings>=

```
assembly component
  add    is add
  {"mrb","rmb"} is b
```

The first rule maps add to itself and the second maps any string that matches mrb or rmb to b. The least general rule that matches a string is applied.

It is often useful to define a default mapping, i.e., for the pattern ``\*'. On the Pentium, for example, most constructor names do not map directly to assembly names, so the default maps a name to the empty string. *This is wrong.*

#### <pentium-linux-default.spec>= (U->)

```
assembly component
  {Indir,{Disp*},Abs32,Reg,{*Index*},E,rel{8,16,32}} is ""
  {*} is $1
```

On the MIPS and SPARC, however, most constructor names map to assembly names, so the default maps a name to itself, i.e., assembly component {\*} is \$1.

The remaining rules specify all the assembly names for the Pentium constructors and illustrate use of the C-shell globbing expressions. In globbing expressions, ``\*' matches any string; any character and ``.' matches itself. The concatenation operator is implicit, so adjacent characters are concatenated. Alternatives are comma-separated lists of strings delimited by ``' and ``'.

We provide one extension to the C-shell syntax: expressions in curly braces may be referenced on the right-hand side by \$*n*, where *n* is the *n*-th braced expression on the left-hand side. For example, the first rule specifies that the suffixes ow and aw map to the assembly name w, and od and ad map to l. The rest of these rules map all the suffixes used in constructor names to their assembly names.

#### <pentium-linux-names.spec>= (U->) [D->]

```
assembly component
  {iAL,AL}          is b
  {iAX,AX}          is w
  {iEAX,eAX}        is l
  {o,a}d            is l
  {o,a}w            is w
  {.I32,.R64,.lsI32,.lsR64} is l
```

```

{.I16,.R32,.lsR32}      is s
.lsI16                  is w
b.*                      is b
{b,w}                   is $1
d                        is l
B.{Eb.Ib,Eb.CL,Eb.Ib,Ev.Ib} is b
B.{Ev.Ib,Ev.CL}         is w
{.STi,.ST.STi,.STi.St}  is ""
P.STi.ST                is P
.{0,NO,B,NB,Z,NZ,BE,NBE,S,NS,P,NP,L,NL,LE,NLE} is $1

```

Many constructor names contain suffixes that are mnemonics for the opcodes they represent. These suffixes are often eliminated in the assembly name. For example, the assembly names for the immediate opcodes ADD.i and OR.i are ADD and OR, respectively. The following rules truncate these suffixes.

[<pentium-linux-names.spec>+=](#) (U->) [<-D->]

```

assembly component
{CALL}.*               is $1
{CALL}l               is $1
CMPXCHG8B              is CMPXCHG
{CMPXCHG,XADD,XCHG,TEST}.Eb.Gb is $1b
CMPSv                  is CMPS
{CMP*}.*              is $1
{DEC,INC}.*           is $1
{DIV}.*               is $1
{*}.st                is $1
{FLD,FSTP}80          is $1t
{FLD,FSTP}.*          is $1
{FILD,FISTP}.*        is $1
{FICOM*}16            is $1s
{FICOM*}32            is $1l
{FILD,FISTP*}64       is $1ll
{IDIV,IMUL}.*         is $1
{IN,INT,J}.*          is $1
JMP.Ep                is lJMP
{JMP}.*               is $1
MOV{.Eb.Ib,.AL.Ob,.Ob.AL} is MOVb
{MOViv,MOV.Ev.Iv}     is MOV
MOVSV.Gv.Ew           is MOVSwl
{MOV.Ew.Sw,MOV.Sw.Ew} is MOVw
{MOVS,MOVZ}X.Gv.Eb    is $1b
{MOVSV,MOVSV.*}       is MOVS
{MOV,MOVS}.*          is $1
{MOVSV,MOVZX}.*       is $1
MOVi{b,w}             is MOV$1
MOVid                 is MOVl
{*}.AX                is $1
{OUT.Ib.AL,OUT.DX.AL} is OUTb
{OUT,OUTS}.*          is $1
{RET.far}*            is lRET
{POP,PUSH,RET}.*      is $1
{SCAS,STOS}v          is $1
{SHRD,SHLD}.*         is $1
SHRSAL                is SHR
TEST{.*.Ib,.Eb.*}     is TESTb
TEST.*.Iw             is TESTw
TEST.*.Id             is TESTl
TEST.*                is TEST
{XADD*}.*             is $1
{XCHG*}.*             is $1
{*}i                  is $1

```

The remaining rules are for constructors with special assembly names.

[<pentium-linux-names.spec>+=](#) (U->) [<-D->]

assembly component

{"mr", "rm"}	is ""
{"mrb", "rmb"}	is b
{*}64	is \$1
{IDIV, DIV}"AL"	is \$1
{IDIV, DIV}"AX"	is \$1
{IDIV, DIV}"eAX"	is \$1
{IMULrm}	is IMUL
INT3	is INT
FLD.ext	is FLDLL
{Jv, Jb}	is J
{INC}.Eb	is INCB
{INS, LODS}v	is \$1
MUL.AL	is MULb
OUTSv	is OUTS
SHLSAL	is SHL
TEST.Ew.Iw	is TESTw
TEST.Ed.Id	is TESTl
SETb	is SET

Component-wise mapping doesn't work for all names.

[<pentium-linux-names.spec>+=](#) (U->) [[<-D->](#)]

```
assembly opcode
CALL.{Ev}{od}          is CALL
CALL.{Jv,Ep}{od,ow}    is lCALL
CALL.aP{od}            is CALL
CMPSv{od,ow}ad         is CMPSl
CMPSv{od,ow}aw         is CMPSw
JMP.Epod               is lJMP
MOVSX.Gv.Ebod          is MOVSbl
MOVSX.Gv.Ebow          is MOVSbw
{ROL,ROR,RCL,RCR,SHR,SAR}{B.Ev.*}od is $1l
{ROL,ROR,RCL,RCR,SHR,SAR}{B.Ev.*}ow is $1w
SHLSAL{B.Ev.*}od       is SHLl
SHLSAL{B.Ev.*}ow       is SHLw
XCHGeAXow              is XCHGw
XCHGeAXod              is XCHGl
```

### Assembly formats for operands

assembly operand introduces mappings from operands to formatted strings that specify how to print the operands in assembly code. Operands may be fields, integer inputs, relocatable addresses, or constant strings. We use printf-style syntax for formatted strings.

The first assembly operand rule below specifies that immediate operands are prefixed by ``\$" and are printed as integers. The second rule specifies that the listed field inputs are prefixed by ``%" and printed as strings, using the names provided in their fieldinfo declarations.

[<pentium-linux-names.spec>+=](#) (U->) [[<-D->](#)]

```
assembly operand
[count i8 i16 i32]      is "$%d"
[r32 sr16 r16 r8 base index] is "%%s"
```

Some inputs are not declared as fields but should be printed in the same format as fields. For example, the input reg should be printed using the names associated with the field base. The following rule uses the optional using *field* clause to specify that reg, reg8, etc. should be printed using the fieldinfo associated with base.

[<pentium-linux-names.spec>+=](#) (U->) [[<-D->](#)]

```
assembly operand
[reg reg8 sreg cr dr]    is "%%s" using field base
```

Some operands are used implicitly in a constructor. For example, the constructor OUT.DX.AL "dx", "al", implicitly uses the dx and al registers as its operands. Like all other operands, their format is assembler dependent so we provide mappings for printing them in the Gnu-Linux format.

[<pentium-linux-names.spec>+=](#) (U->) [[<-D](#)]

```
assembly operand
```



```
dx    is  "%dx"
ax    is  "%ax"
```

## Assembly syntax for constructors

The default assembly syntax for a constructor appears in the constructor's specification; an alternate syntax may be specified with `assembly syntax`. Providing assembly syntax in a constructor specification can help a specification writer or user read and identify a constructor, and it is concise when only one assembly syntax is required. An alternate syntax may be needed, however, if more than one assembly language is used on the target.

`assembly syntax` uses the same syntax as the constructors directive: a constructor name followed by a list of operands. The assembly-syntax specification must use the same set of operands that the constructor uses, but the operands may appear in any order and with any syntactic sugar. To reduce redundancy, we define new patterns to group constructors that share the same assembly syntax.

The Gnu-Linux assembly language reverses the order of all operands. The following directives reverse the order.

[`<pentium-linux-syntax.spec>= \(U->\) \[D->\]`](#)

assembly syntax

```
arith^"iAL"      i8!, "%al"
arith^"iAX"      i16!, "%ax"
arith^"iEAX"     i32!, "%eax"
DIV^"AL"        Eaddr, "%al"
DIV^"AX"        Eaddr, "%ax"
DIV^"eAX"       Eaddr, "%eax"

arithI^"b"       i8!, Eaddr
arithI^"w"       i16!, Eaddr
arithI^"d"       i32!, Eaddr
arithI^"ov^"b"   i8!, Eaddr
MOV.Eb.Ib       i8!, Eaddr
MOV.Ev.Iv^ow    i16!, Eaddr
MOV.Ev.Iv^od    i32!, Eaddr
```

[`<pentium-linux-syntax.spec>+= \(U->\) \[<-D->\]`](#)

assembly syntax

```
arith^"rmb"      Eaddr, reg8
arith^"rm"^ov    Eaddr, reg
IMULrm^ov       Eaddr, reg
MOV^"rmb"       Eaddr, reg
MOV^"rm"^ov     Eaddr, reg
MOVZX.Gv.Ew     Eaddr, r16
MOVSX.Gv.Ew     Eaddr, r16
MOVZX.Gv.Eb^ov  Eaddr, r32
MOVSX.Gv.Eb^ov  Eaddr, r32
BSF^ov         Eaddr, reg
BSR^ov         Eaddr, reg
LAR^ov         Eaddr, reg

arith^"mr^b"    reg8, Eaddr
arith^"mr"^ov   reg, Eaddr
MOV^"mr"^ov     reg, Eaddr
MOV^"mr^b"     reg, Eaddr
TEST.Ev.Gv^ov  reg, Eaddr
BT^ov          reg, Eaddr
BTi^ov         i8!, Eaddr
BTC^ov         reg, Eaddr
BTCi^ov        i8!, Eaddr
BTR^ov         reg, Eaddr
BTRi^ov        i8!, Eaddr
BTS^ov         reg, Eaddr
BTSi^ov        i8!, Eaddr
CMPXCHG.Eb.Gb  reg, Eaddr
CMPXCHG.Ev.Gv^ov reg, Eaddr
```

[`<pentium-linux-syntax.spec>+= \(U->\) \[<-D->\]`](#)

```

patterns
  fstack    is FADD | FDIV | FDIVR | FMUL | FSUB | FSUBR
  fsti      is fstack | FCOMs
  stidx     is FFREE | FUCOMs | FXCH
  Sstack    is P.STi.ST | .STi.St

```

```

assembly syntax
  fstack^Sstack "%st", "%st"(idx)
  fstack^.ST.STi "%st"(idx), "%st"
  FCOMs^.ST.STi "%st"(idx), "%st"
  FSTs.st^.STi "%st"(idx)
  FLD.STi "%st"(idx)
  stidx "%st"(idx)

```

```

FNSTSW.AX "%ax"
FSTSW.AX "%ax"

```

```

IDIV^"AX" Eaddr, "%ax"
IDIV^"eAX" Eaddr, "%eax"

```

```

IN.AL.Ib i8!, "%al", i8!
IN.eAX.Ib^ov i8!, "%eax", i8!
IN.AL.DX "%dx, %al"
IN.eAX.DX^ov "%dx, %eax"

```

```

IMUL.Iv^"d" i32!, Eaddr, reg
INT3 "$3"
LEA^ov Mem, reg

```

```

MOVib i8!, r8
MOViw i16!, r32
MOVid i32!, r32

```

```

MOV.AL.Ob offset, "%al"
MOV.eAX.Ov^ov offset, "%eax"
MOV.Ob.AL "%al", offset
MOV.Ov.eAX^ov "%eax", offset

```

```

OUT.Ib.AL "%al", i8!
OUT.Ib.eAX^ov "%eax", i8!
OUT.DX.AL "%al", "%dx"
OUT.DX.eAX^ow "%al", "%dx"
OUT.DX.eAX^od "%eax", "%dx"

```

```

patterns
  pES is POP.ES | PUSH.ES
  pSS is POP.SS | PUSH.SS
  pDS is POP.DS | PUSH.DS
  pFS is POP.FS | PUSH.FS
  pGS is POP.GS | PUSH.GS

```

```

assembly syntax
  pES "%ES"
  pSS "%SS"
  pDS "%DS"
  pFS "%FS"
  pGS "%GS"
  PUSH.CS "%CS"

```

```

rot^B.Eb.1 "$1", Eaddr
rot^B.Ev.1^ov "$1", Eaddr

```

```

rot^B.Eb.CL "%cl", Eaddr
rot^B.Ev.CL^ov "%cl", Eaddr

```

```

rot^B.Eb.Ib i8!, Eaddr
rot^B.Ev.Ib^ov i8!, Eaddr

```

```
shdIb^ov count, reg, Eaddr
shdCL^ov "%cl", reg, Eaddr
```

[<pentium-linux-syntax.spec>+=](#) ([U->](#)) [[<-D->](#)]

```
TEST.AL.Ib      i8, "%al"
TEST.eAX.Iv^ow  i16, "%ax"
TEST.eAX.Iv^od  i32, "%ax"
TEST.Eb.Ib      i8, Eaddr
TEST.Ew.Iw      i16, Eaddr
TEST.Ed.Id      i32, Eaddr
TEST.Eb.Gb      reg, Eaddr
TEST.Ev.Gv^ov   reg, Eaddr
XADD.Eb.Gb      reg, Eaddr
XADD.Ev.Gv^ov   reg, Eaddr
XCHG.Eb.Gb      reg, Eaddr
XCHG^"eAX"^ov  "%eax", r32
XCHG.Ev.Gv^ov   reg, Eaddr
```

Effective addresses also use a different syntax under Gnu-Linux.

[<pentium-linux-syntax.spec>+=](#) ([U->](#)) [[<-D](#)]

assembly syntax

```
Indir      (reg)
Disp32     d(reg)
Index      (base,index,ss)
Index32     d(base,index,ss)
ShortIndex d(,index,ss)
```

## Miscellaneous

The Intel 486 instruction set is a subset of the Pentium set. When generating instructions for that target, the pentium-only instructions are discarded.

[<i486-linux.spec>=](#)

[<pentium32.spec>](#)

discard [<names of pentium constructors>](#)

[<pentium-linux.spec>](#)

Names for generating assembly emitters are specified in several chunks.

[<pentium-names.spec>=](#)

[<pentium-linux-names.spec>](#)

[<pentium-linux-default.spec>](#)

[<pentium-linux-syntax.spec>](#)

[<header.spec>=](#) ([U->](#) [U->](#) [U->](#) [U->](#) [U->](#) [U->](#) [U->](#))

[<field specs>](#)

patterns [<patterns for integer opcodes>](#)

[<pattern specs for other patterns>](#)

[<prefix assignments>](#)

[<placeholders>](#)

[<constructors for displacements>](#)

[<constructors for effective addresses>](#)

Check the arithmetic constructors. Substitute other chunks for this one to check each chunk.

[<pentium-AF.spec>=](#)

[<header.spec>](#)

[<alphabetical constructors A-F>](#)

[<pentium-GK.spec>=](#)

[<header.spec>](#)

[<alphabetical constructors G-K>](#)

[<pentium-LQ.spec>=](#)  
[<header.spec>](#)  
[<alphabetical constructors L-Q>](#)

[<pentium-R.spec>=](#)  
[<header.spec>](#)  
[<alphabetical constructors R>](#)

[<pentium-SZ.spec>=](#)  
[<header.spec>](#)  
[<alphabetical constructors S-Z>](#)

[<pentium-arith.spec>=](#)  
[<header.spec>](#)  
[<arithmetic constructors>](#)

[<pentium-float.spec>=](#)  
[<header.spec>](#)  
[<floating-point constructors>](#)

The Pentium checker checks all 32-bit instructions accepted by the Gnu-Linux assembler. We omit some variants of the rot instructions, because exhaustively checking each one seemed unnecessary.

[<pentium-check.spec>=](#)  
[<pentium32.spec>](#)  
[<pentium-linux.spec>](#)  
discard  
rot^B.Eb.1 rot^B.Ev.CL ROL^B.Eb.Ib RCL^B.Eb.Ib  
SHLSAL^B.Eb.Ib SAR^B.Eb.Ib ROR^B.Ev.Ib^ov  
RCR^B.Ev.Ib^ov SHR^B.Ev.Ib^ov

[<pentium-checker.s>=](#)  
.align 16

- [<alphabetical constructors>](#): [U1](#), [D2](#)
- [<alphabetical constructors A-F>](#): [U1](#), [D2](#), [D3](#), [D4](#), [D5](#), [D6](#), [D7](#), [D8](#), [D9](#), [D10](#), [D11](#), [D12](#), [D13](#), [D14](#), [D15](#), [D16](#), [D17](#), [D18](#), [D19](#), [D20](#), [D21](#), [D22](#), [D23](#), [D24](#), [D25](#), [D26](#), [D27](#), [D28](#), [D29](#), [D30](#), [D31](#), [U32](#)
- [<alphabetical constructors G-K>](#): [U1](#), [D2](#), [D3](#), [D4](#), [D5](#), [U6](#)
- [<alphabetical constructors L-Q>](#): [U1](#), [D2](#), [D3](#), [D4](#), [D5](#), [D6](#), [D7](#), [D8](#), [D9](#), [U10](#)
- [<alphabetical constructors R>](#): [U1](#), [D2](#), [D3](#), [U4](#)
- [<alphabetical constructors S-Z>](#): [U1](#), [D2](#), [D3](#), [D4](#), [D5](#), [D6](#), [D7](#), [U8](#)
- [<arithmetic constructors>](#): [U1](#), [D2](#), [U3](#)
- [<32-bit constructors>](#): [U1](#), [D2](#), [D3](#), [D4](#), [D5](#), [D6](#), [D7](#), [D8](#), [D9](#), [D10](#), [D11](#), [D12](#), [D13](#), [D14](#), [D15](#), [D16](#), [D17](#), [D18](#), [D19](#), [D20](#), [D21](#), [D22](#), [D23](#), [D24](#), [D25](#)
- [<constructors for displacements>](#): [U1](#), [D2](#), [U3](#)
- [<constructors for effective addresses>](#): [U1](#), [D2](#), [U3](#)
- [<example mappings>](#): [D1](#)
- [<field specs>](#): [U1](#), [D2](#), [D3](#), [D4](#), [D5](#), [D6](#), [D7](#), [D8](#), [D9](#), [U10](#)
- [<floating-point constructors>](#): [U1](#), [U2](#), [D3](#), [D4](#), [D5](#), [D6](#), [D7](#), [D8](#), [D9](#), [D10](#), [D11](#), [D12](#), [D13](#), [D14](#), [D15](#), [D16](#), [D17](#), [D18](#), [D19](#), [D20](#), [D21](#), [D22](#), [D23](#), [U24](#)
- [<header.spec>](#): [D1](#), [U2](#), [U3](#), [U4](#), [U5](#), [U6](#), [U7](#), [U8](#)
- [<i486-linux.spec>](#): [D1](#)
- [<mld-pentium.spec>](#): [D1](#), [D2](#)
- [<more fields of opcode>](#): [U1](#), [D2](#), [D3](#), [D4](#), [D5](#)
- [<names of pentium constructors>](#): [D1](#), [D2](#), [D3](#), [D4](#), [D5](#), [D6](#), [U7](#)
- [<pattern specs for other patterns>](#): [U1](#), [D2](#), [D3](#), [D4](#), [D5](#), [D6](#), [D7](#), [D8](#), [D9](#), [D10](#), [D11](#), [D12](#), [D13](#), [U14](#)
- [<patterns for integer opcodes>](#): [U1](#), [D2](#), [D3](#), [D4](#), [D5](#), [D6](#), [D7](#), [D8](#), [D9](#), [D10](#), [D11](#), [D12](#), [D13](#), [D14](#), [D15](#), [D16](#), [D17](#), [D18](#), [D19](#), [D20](#), [D21](#), [D22](#), [U23](#)
- [<pentium-AF.spec>](#): [D1](#)
- [<pentium-arith.spec>](#): [D1](#)
- [<pentium-checker.s>](#): [D1](#)
- [<pentium-check.spec>](#): [D1](#)
- [<pentium-core.spec>](#): [D1](#)
- [<pentium-float.spec>](#): [D1](#)
- [<pentium-GK.spec>](#): [D1](#)

- [<pentium-int.spec>](#): [D1](#)
- [<pentium-linux-default.spec>](#): [D1](#), [U2](#)
- [<pentium-linux-names.spec>](#): [D1](#), [D2](#), [D3](#), [D4](#), [D5](#), [D6](#), [D7](#), [U8](#)
- [<pentium-linux.spec>](#): [D1](#), [D2](#), [D3](#), [D4](#), [D5](#), [D6](#), [U7](#), [U8](#)
- [<pentium-linux-syntax.spec>](#): [D1](#), [D2](#), [D3](#), [D4](#), [D5](#), [U6](#)
- [<pentium-LQ.spec>](#): [D1](#)
- [<pentium-names.spec>](#): [D1](#)
- [<pentium-R.spec>](#): [D1](#)
- [<pentium32.spec>](#): [D1](#), [U2](#), [U3](#)
- [<pentium-synth.spec>](#): [D1](#)
- [<pentium-SZ.spec>](#): [D1](#)
- [<placeholders>](#): [U1](#), [D2](#), [D3](#), [D4](#), [U5](#)
- [<prefix assignments>](#): [U1](#), [D2](#), [D3](#), [U4](#)
- [<proposed constructors for effective addresses>](#): [D1](#)
- [<synthetics with WAIT>](#): [D1](#), [D2](#), [D3](#), [D4](#), [D5](#), [D6](#), [U7](#)